

# Trace Based Just-In-Time Type Specialization for Dynamic Languages ( PLDI 2009 )

**Gal, Eich, Shaver, Anderson, Mandelin, Kaplan, Hoare, Zbarsky, Orendorff,  
Ruderman, Smith, Reitmaier, Haghighat, Bebenita, Chang, Franz**

16th February 2015

# Outline

# Motivation

- Generate efficient machine code for dynamic type languages like Javascript.
- Challenges
  - Types decided at runtime
  - Even type inferencing does not help much

# Motivation

- Generate efficient machine code for dynamic type languages like Javascript.
- Challenges
  - Types decided at runtime
  - Even type inferencing does not help much

# Motivation

- Generate efficient machine code for dynamic type languages like Javascript.
- Challenges
  - Types decided at runtime
  - Even type inferencing does not help much

# TraceMonkey: Proposed Compilation strategy

- Is implemented for Javascript Interpreter SpiderMonkey
- Dynamic compilation on loop level granularity, not method based. Why??

# TraceMonkey: Proposed Compilation strategy

- Is implemented for Javascript Interpreter SpiderMonkey
- Dynamic compilation on loop level granularity, not method based. Why??

# Trace

- Typed loop traces
  - Single entry, multiple exits loop paths
  - Type specialized compilation



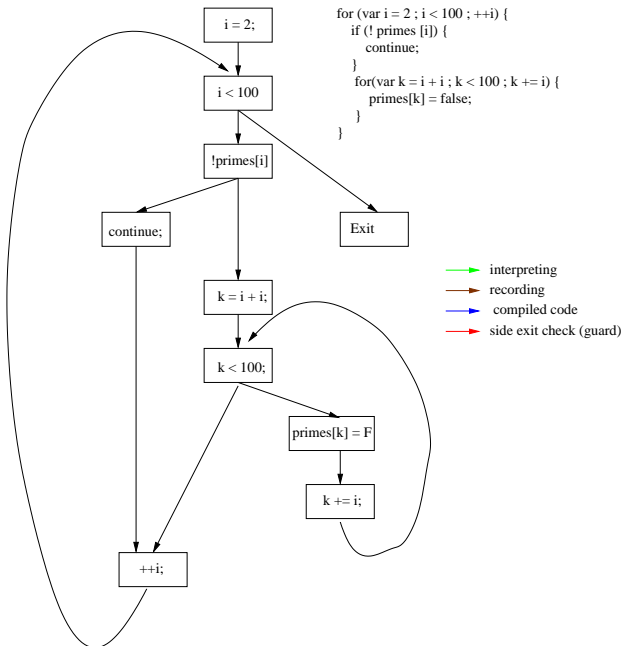
# Trace

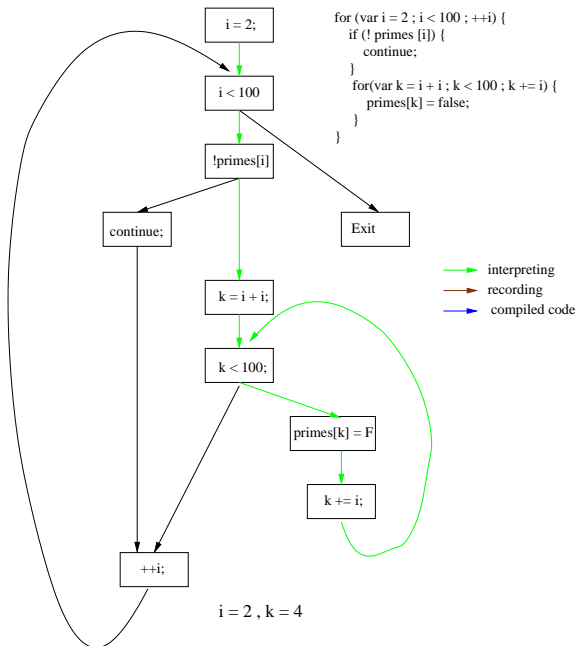
- Typed loop traces
  - Single entry, multiple exits loop paths
  - Type specialized compilation

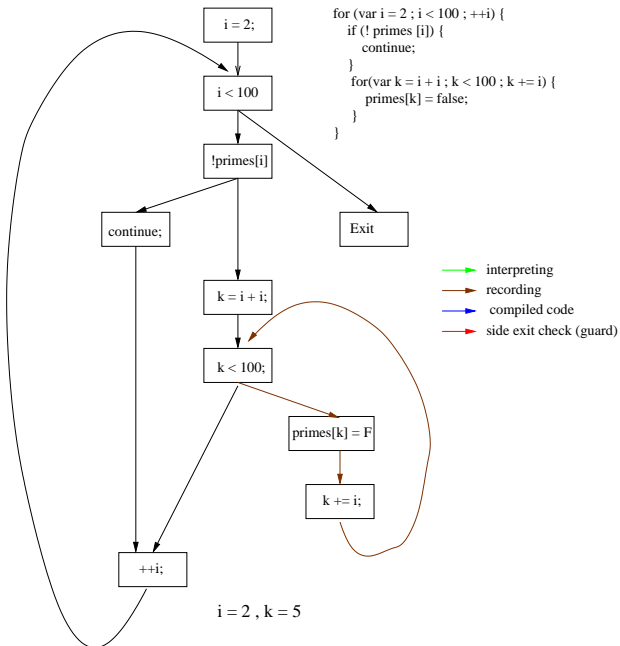
# Trace

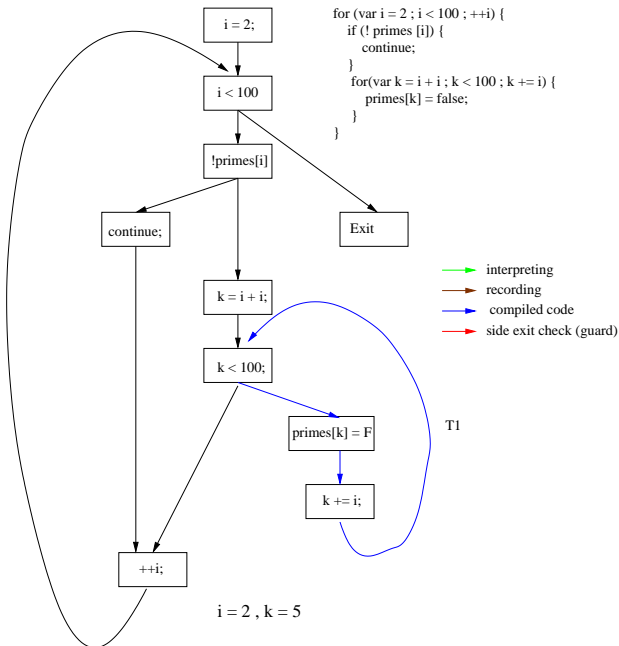
- Typed loop traces
  - Single entry, multiple exits loop paths
  - Type specialized compilation

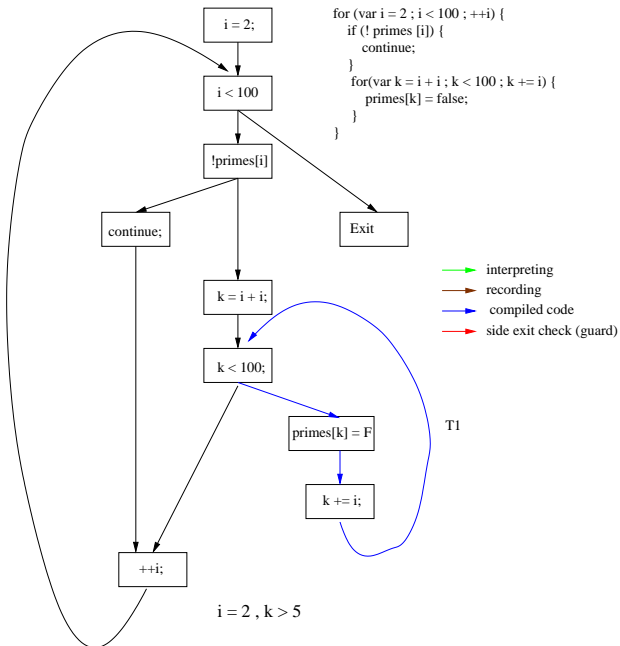
# Outline



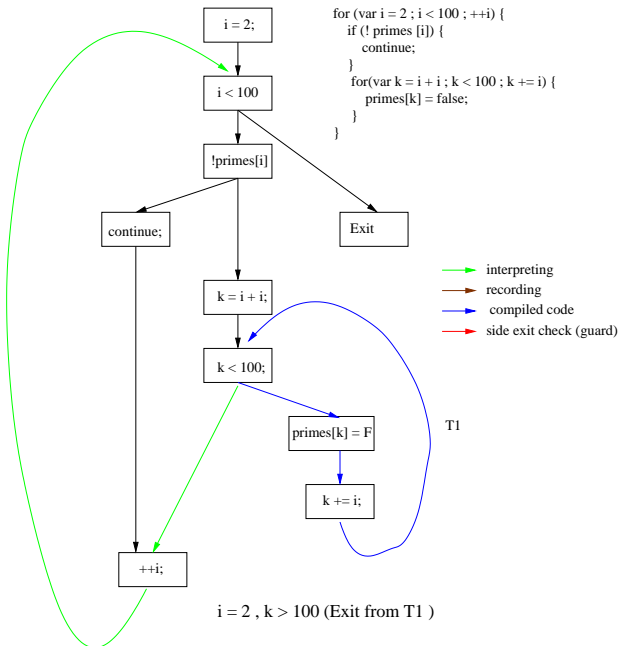


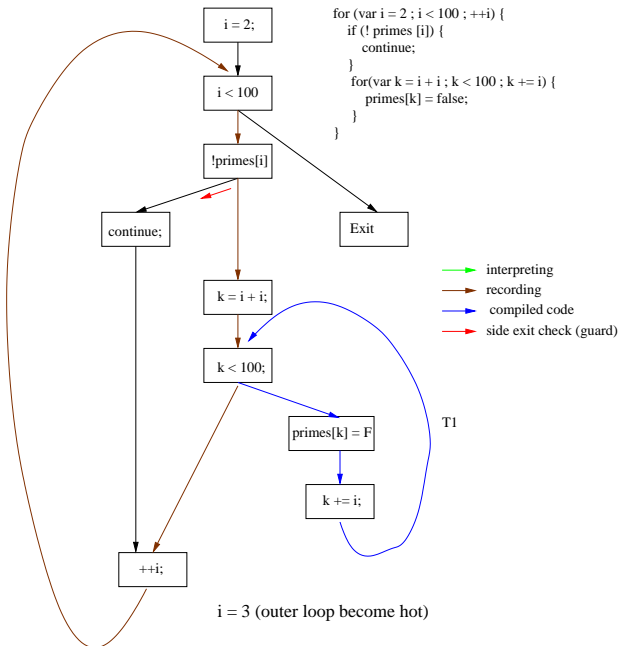


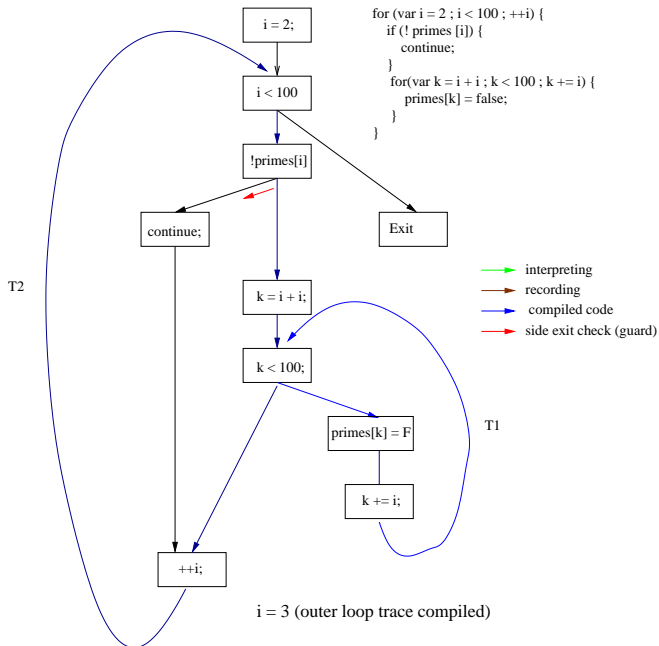


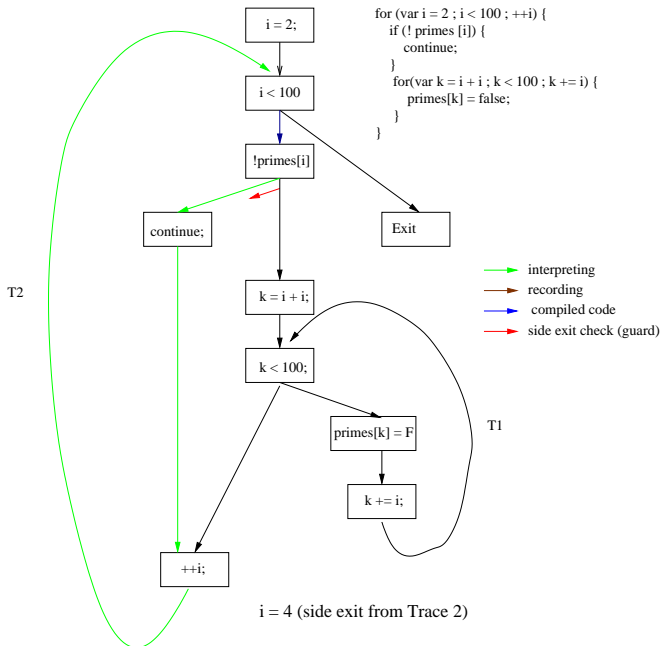


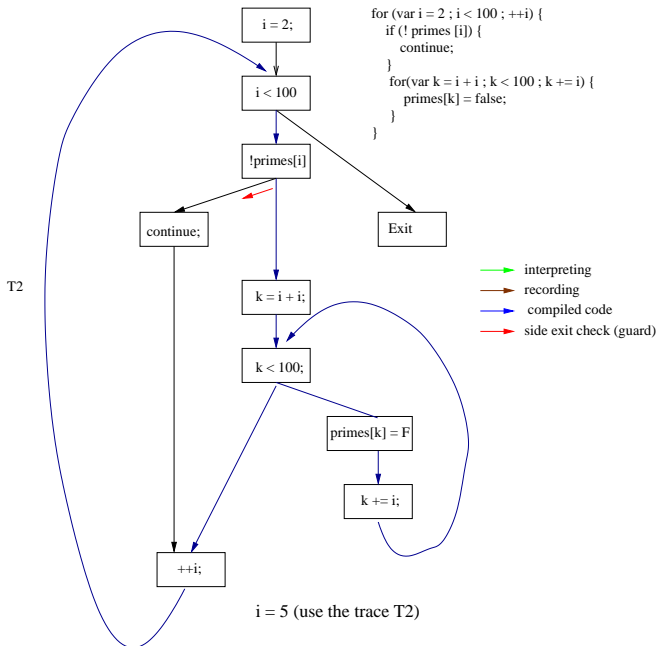


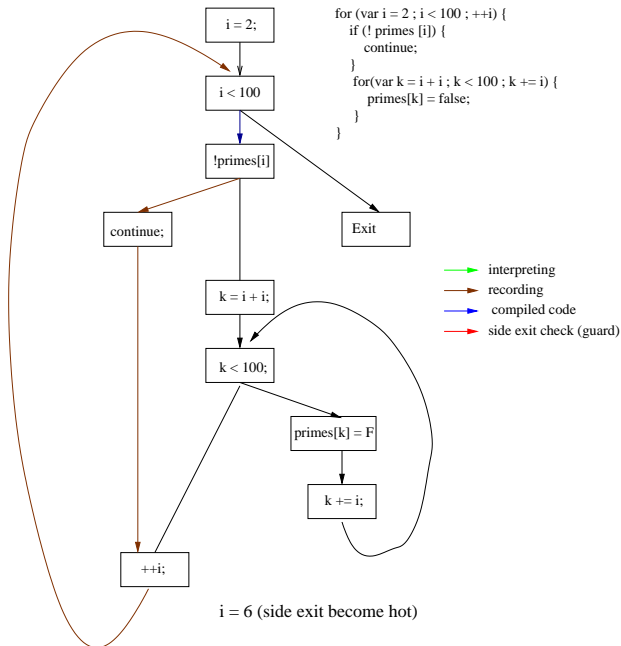


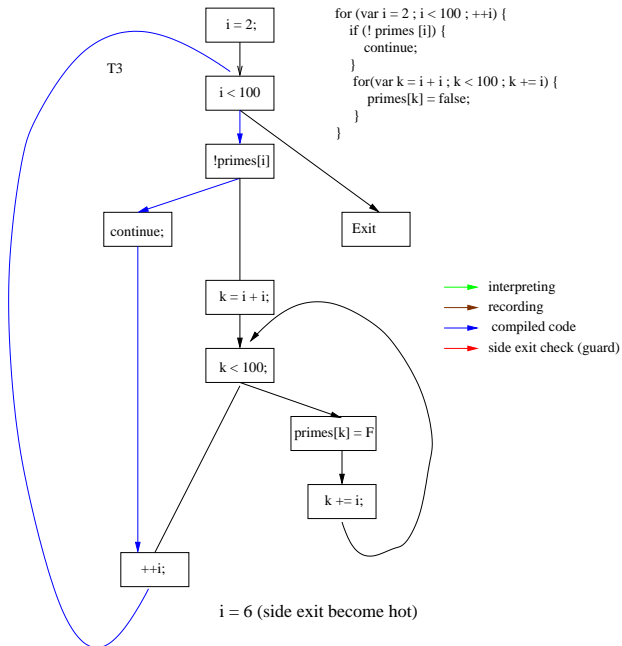














# Outline




# Type Stable Trace




# Type Stable Trace



# Type Stable Trace




# Trace Tree Extension



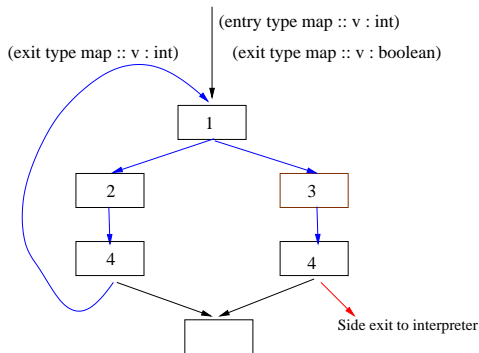
Note: Current Implementation extends only if

- side exit is for control flow branch
- does not leave the loop

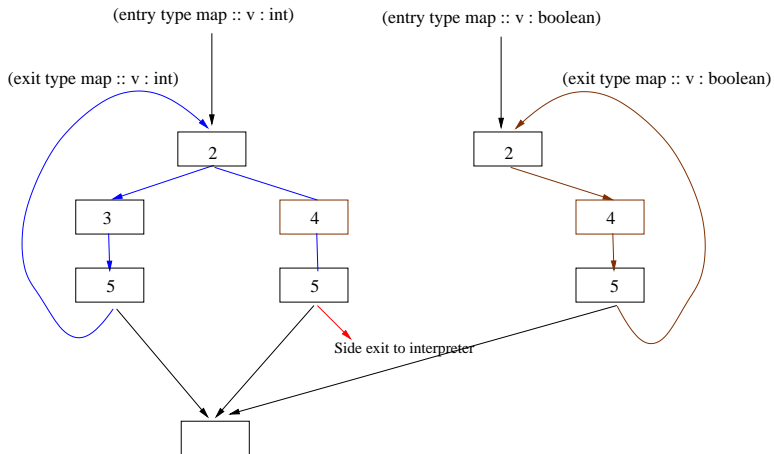
# Type Unstable Trace Handling



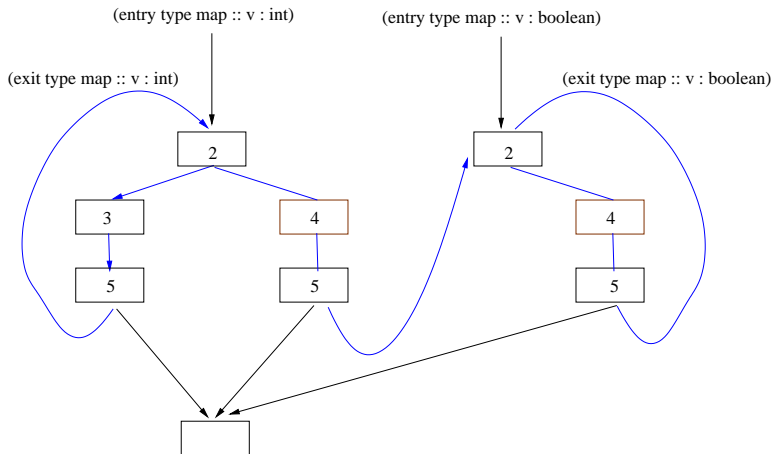
# Type Unstable Trace Handling



# Type Unstable Trace Handling



# Type Unstable Trace Handling





# Blacklisting


- What if a hot loop contain traces that always fail recording ?
- Solution : Blacklist with backoff on recording

# Blacklisting


- What if a hot loop contain traces that always fail recording ?
- Solution : Blacklist with backoff on recording

# Outline


# Nested Trace Tree Formation




# Nested Trace Tree Formation




# Nested Trace Tree Formation



# Nested Trace Tree Formation





# Nested Trace Tree Formation






# Continue tracing for outer loop

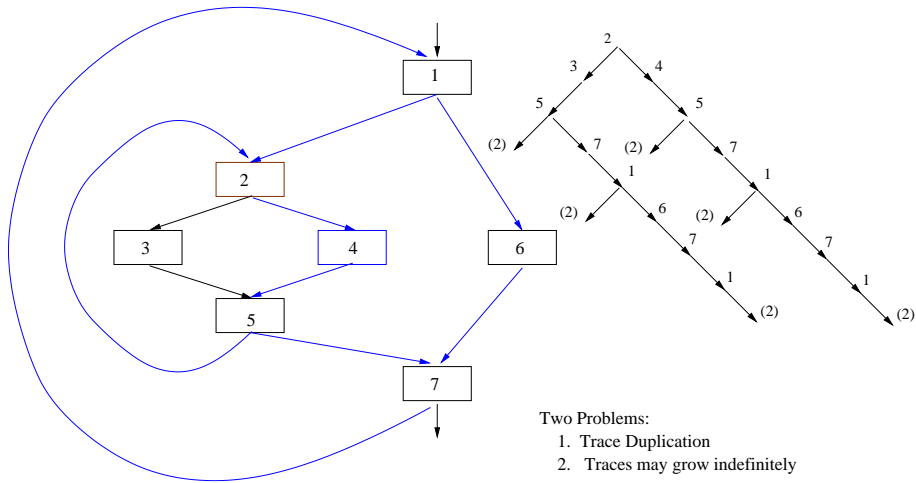





# Continue tracing for outer loop



Continue tracing for outer loop




# Separate traces




Outer loop starts recording

# Separate traces




Outer Loop recording stops as "type matched" inner compiled trace not available.

# Separate traces



Inner compiled trace ready


# Separate traces



Outer loop records the inner trace call in its recording




# Separate traces




What if during recording of outer loop, the inner trace got a side exit !!

# Separate traces



# Separate traces



What if during execution of outer loop, the inner trace got a side exit !!

# Modified Blacklisting with Nesting

- The outer loop blacklisted very quickly.
- Solution : Blacklist with an option of “forgiving”

# Modified Blacklisting with Nesting

- The outer loop blacklisted very quickly.
- Solution : Blacklist with an option of “forgiving”

# Outline

# Recording

- While interpreting traces are recorded and transformed into LIR. WHY ??
  - LIR traces are type specialized
  - LIR traces are representation specialized w.r.t objects
  - LIR traces are representation specialized w.r.t numbers
  - LIR traces inline functions
- All the above LIR features need guards

# Recording

- While interpreting traces are recorded and transformed into LIR. WHY ??
  - LIR traces are type specialized
  - LIR traces are representation specialized w.r.t objects
  - LIR traces are representation specialized w.r.t numbers
  - LIR traces inline functions
- All the above LIR features need guards



# Recording

- While interpreting traces are recorded and transformed into LIR. WHY ??
  - LIR traces are type specialized
  - LIR traces are representation specialized w.r.t objects
  - LIR traces are representation specialized w.r.t numbers
  - LIR traces inline functions
- All the above LIR features need guards

# Recording

- While interpreting traces are recorded and transformed into LIR. WHY ??
  - LIR traces are type specialized
  - LIR traces are representation specialized w.r.t objects
  - LIR traces are representation specialized w.r.t numbers
  - LIR traces inline functions
- All the above LIR features need guards

# Recording

- While interpreting traces are recorded and transformed into LIR. WHY ??
  - LIR traces are type specialized
  - LIR traces are representation specialized w.r.t objects
  - LIR traces are representation specialized w.r.t numbers
  - LIR traces inline functions
- All the above LIR features need guards

# Recording

- While interpreting traces are recorded and transformed into LIR. WHY ??
  - LIR traces are type specialized
  - LIR traces are representation specialized w.r.t objects
  - LIR traces are representation specialized w.r.t numbers
  - LIR traces inline functions
- All the above LIR features need guards

# Trace Tree Optimization

- Optimization Goal: Make compilation fast. (Done by NanoJIT)
  - Restrict to small set of optimization
  - Avoid any reordering optimizations
  - leverage the fact that the LIR is in SSA form
- Optimizations are performed in 2 phases: forward and backward

# Trace Tree Optimization

- Optimization Goal: Make compilation fast. (Done by NanoJIT)
  - Restrict to small set of optimization
    - Avoid any reordering optimizations
    - leverage the fact that the LIR is in SSA form
- Optimizations are performed in 2 phases: forward and backward

# Trace Tree Optimization

- Optimization Goal: Make compilation fast. (Done by NanoJIT)
  - Restrict to small set of optimization
  - Avoid any reordering optimizations
  - leverage the fact that the LIR is in SSA form
- Optimizations are performed in 2 phases: forward and backward

# Trace Tree Optimization

- Optimization Goal: Make compilation fast. (Done by NanoJIT)
  - Restrict to small set of optimization
  - Avoid any reordering optimizations
  - leverage the fact that the LIR is in SSA form
- Optimizations are performed in 2 phases: forward and backward



# Trace Tree Optimization

- Optimization Goal: Make compilation fast. (Done by NanoJIT)
  - Restrict to small set of optimization
  - Avoid any reordering optimizations
  - leverage the fact that the LIR is in SSA form
- Optimizations are performed in 2 phases: forward and backward

# Forward Optimizations

- While recording on each instruction basis.
  - Constant subexpression elimination
  - Expression simplification and algebraic identities
  - ISA specific simplification

# Forward Optimizations

- While recording on each instruction basis.
  - Constant subexpression elimination
  - Expression simplification and algebraic identities
  - ISA specific simplification

# Forward Optimizations

- While recording on each instruction basis.
  - Constant subexpression elimination
  - Expression simplification and algebraic identities
  - ISA specific simplification

# Forward Optimizations

- While recording on each instruction basis.
  - Constant subexpression elimination
  - Expression simplification and algebraic identities
  - ISA specific simplification

# Backward optimizations

- When recording is complete.
  - dead data-stack store elimination
  - dead call-stack store elimination
  - dead code elimination
  - Register Allocation

# Backward optimizations

- When recording is complete.
  - dead data-stack store elimination
  - dead call-stack store elimination
  - dead code elimination
  - Register Allocation

# Backward optimizations

- When recording is complete.
  - dead data-stack store elimination
  - **dead call-stack store elimination**
  - dead code elimination
  - Register Allocation



# Backward optimizations

- When recording is complete.
  - dead data-stack store elimination
  - dead call-stack store elimination
  - **dead code elimination**
  - Register Allocation

# Backward optimizations

- When recording is complete.
  - dead data-stack store elimination
  - dead call-stack store elimination
  - dead code elimination
  - Register Allocation

# Outline

# Trace stitching

- when a trace calls a branch trace, trace stitching can be applied
- identical type maps yield identical activation records

# Outline

# Setup

- SunSpider Benchmark suite is used
- SunSpider test driver starts a JS interpreter, load and run each program
  - one time to warm up
  - 10 times, which are measured the average time is taken
  - runs on 4 interpreter
    - SpiderMonkey: JS interpreter. **Baseline for comparison**
    - TraceMonkey: The proposed compilation strategy
    - SquirrelFish Extreme (SFX): Call threaded JS interpreter
    - V8: Method compiling JS VM

# Results

- Speed up spreads between 0.9 and 25.
- Performance breakup on Sunspider benchmark
  - 9 / 26 : TraceMonkey is the fastest
  - 5 / 26 : The current implementation does not trace recursion.
  - 5 / 26 : nested loop with small bodies. May be more time in calling nested trace.
  - 2 / 26 : Does not trace eval and some other C implemented functions.
  - 2 / 26 : Trace well, but long compilation.
  - 1 / 26 : Does not trace through regular expression replace operations.
  - 1 / 26 : run time dominated by string processing builtins.
  - 1 / 26 : dominated by regular expression matching (implemented in all 3 VM as special regular expression compiler)

# Results Continued

- Overall performance speedup of native trace execution = 
$$\frac{\text{Time per bytecode execution in Interpreter}}{\text{Time per bytecode execution in native code}} = 3.9$$
- Other interpreter have an overall performance speedup of 3.0
- TraceMonkey recording & compilation 200 times slower than interpreter speed.
- the performance has caught up after 270 iterations of a trace.



# My Critique

- Web as Desktop
- Incremental recompilation.

# Outline



# Outline

# Correctness

- Mozillas JavaScript fuzz tester, JSFUNFUZZ, random test generator.
- Modified JSFUNFUZZ to generate
  - loops
  - type unstable loops
  - heavy branching code.

# Implementation

- Implemented for SpiderMonkey JS virtual machine.
- SpiderMonkey is the interpreter for JS
- first compiled into bytecode, then interpreted.
- is garbage collected (non-generational, stop the world, mark and sweep)

# Calling compiled trace

- traces are stored in a trace cache indexed by interpreter PC and type map.
- traces are compiled so that they can be called as functions using standard native calling conventions.