

# Sample Sort Using MPI

---

## Experimentation data for Serial computation

For serial computation we have chosen the serial code with 12 buckets and obtained the runtimes with size = 120000, 1200000, 12000000, 120000000. These results will be used as the baseline to compare with the parallel models.

(N)	Runtime (sec)
120000	0.058558
1200000	0.407464
12000000	3.881106
120000000	44.601656

## MPI Implementation

### Overview of the Algorithm and suggested optimization

1. Rank 0 creates the data set and scatters it to the P processes.
2. Each process does local sorting.
3. Each process select p-1 splitters.
4. All the p(p-1) splitters are gathered to rank 0.
5. Rank 0 selects a set of p-1 global splitters and sorting the sample of p(p-1) splitters and broadcasted that to P processes.
6. Each process put its local data set into local buckets based on the splitters set.
7. Each process rearranges its buckets using MPI\_Alltoall so that P0 has all the data < splitter[0] and soon.
8. Each process finally sorts their local buckets.



### Experimentation data

#### Tabular Representation

**Note:** The speedup is calculated using (serial runtime calculated above / runtime for P).

P	Runtime (seconds)	Speedup
N - 120000		
2	.032027	1.828349
6	.010830	5.407017
12	.021130	2.771294
N - 1200000		
2	.258872	1.573996
6	.1213530	3.357675
12	.074790	5.4480497
N - 12000000		
2	2.206225	1.759161
6	.796015	4.875664
12	.499257	7.773751
N - 120000000		
2	23.916551	1.864886
6	9.023845	4.942643
12	5.772662	7.726357

## Graphical Representation



### Analysis

**Q.** Discuss the results you obtained. Were you able to gain a linear or near-linear speedup? Why or why not?

**Answer.**

### Discussion on results

1. For  $N = 120000$ , the speedup drops down after  $P = 6$ .
  - Here the communication cost involved when  $P > 6$  is much greater than the computation gain by sharing the workload among processors and as a result the performance degrades after  $P = 6$ .
2. For a given  $N$ , the speedup increases with value of  $P$ .
  - With more number of  $P$ , the workload per process ( $N/P$ ) decreases leading to better overall performance (or lesser runtime).
3. With very large values of  $P$ , the performance curve breaks the linear behavior.
  - Let take the example of  $N=1200000$ . Had the curve been linear, then the expected value of speedup at  $P = 12$  is 6.71535, but we got 5.448. Similarly, for  $N = 12000000$ , the obtained value of speedup is 7.73 at  $P = 12$  which degrades from the expected value of 9.75 (which should be if the curve is linear). The same is true for  $N = 120000000$ . In the examples mentioned above, the communication cost involved when  $P = 12$  is much greater than the computation gain by sharing the workload among processors and as a result the performance degrades at  $P = 12$ .

### Linear or near linear speedup

- The speedup that we obtain **is near linear**.
  - With initial incremental values of  $P$ , the speedup is increasing proportionately as the workload is distributed proportionately among the processes. But with larger values of  $P$ , the communication cost outweighs the gain with workload distribution and as a result the speedup curve started degrading at values where the computation/ communication ratios are low. (This is substantiated in Analysis 3 above.)

# Histogram Sort Using MPI

---

## Experimentation data for Serial computation

For serial computation we have chosen the serial code with 12 buckets and obtained the runtimes with size = 120000, 1200000, 12000000, 120000000. These results will be used as the baseline to compare with the parallel models.

(N)	Runtime (sec)
120000	0.060661
1200000	0.449810
12000000	4.006512
120000000	45.562630

## MPI Implementation


### Experimentation data


#### Tabular Representation

**Note:** The speedup is calculated using (serial runtime calculated above / runtime for P).

P	Runtime (seconds)	Speedup
N - 120000		
2	.034180	1.774720
6	.013763	4.407349
12	.034938	1.736207
N - 1200000		
2	.283378	1.587312
6	.127599	3.525173
12	.078755	5.711481
N - 12000000		
2	2.280247	1.757051
6	.873000	4.589360
12	.5860612	6.836337
N - 120000000		
2	25.159230	1.810970
6	9.624384	4.734082
12	6.537176	6.969772

### Graphical Representation





### Analysis

**Q.** Discuss the results you obtained. Were you able to gain a linear or near-linear speedup? Why or why not?

**Answer.**

### Discussion on results

- For  $N = 120000$ , the speedup drops down after  $P = 6$ .
  - Here the communication cost involved when  $P > 6$  is much greater than the computation gain by sharing the workload among processors and as a result the performance degrades after  $P = 6$ .
- For a given  $N$ , the speedup increases with value of  $P$ .
  - With more number of  $P$ , the workload per process ( $N/P$ ) decreases leading to better overall performance (or lesser runtime).
- With very large values of  $P$ , the performance curve breaks the linear behavior.
  - Let take the example of  $N=1200000$ . Had the curve been linear, then the expected value of speedup at  $P = 12$  is 7.05, but we got 5.711481. Similarly, for  $N = 12000000$ , the obtained value of speedup is 6.83 at  $P = 12$  which degrades from the expected value of 9.17 (which should be if the curve is linear). The same is true for  $N = 120000000$ . In the examples mentioned above, the communication cost involved when  $P = 12$  is much greater than the computation gain by sharing the workload among processors and as a result the performance degrades at  $P = 12$ .

### Linear or near linear speedup

- The speedup that we obtain **is near linear**.
  - With initial incremental values of  $P$ , the speedup is increasing proportionately as the workload is distributed proportionately among the processes. But with larger values of  $P$ , the communication cost outweighs the gain with workload distribution and as a result the speedup curve started degrading at values where the computation/ communication ratios are low. (This is substantiated in Analysis 3 above.)